



Internet of Things and Services

Service-oriented architectures


Docker

Container-based virtualization

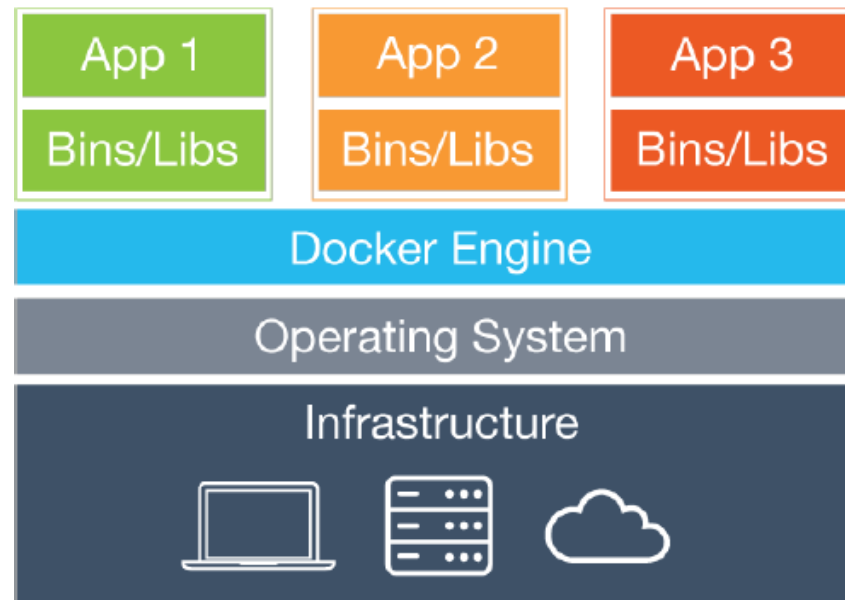
Department of Computer Science
Faculty of Electronic Engineering, University of Nis

Internet of Things and Services
Computing and informatics

Prof. dr Dragan Stojanović



- ❁ Lightweight, open and secure container-based virtualization
 - ❁ Containers include the application and all of its dependencies, but share the OS kernel with other containers
 - ❁ Containers run as an isolated process in user space on the host OS
 - ❁ Containers are also not tied to any specific infrastructure



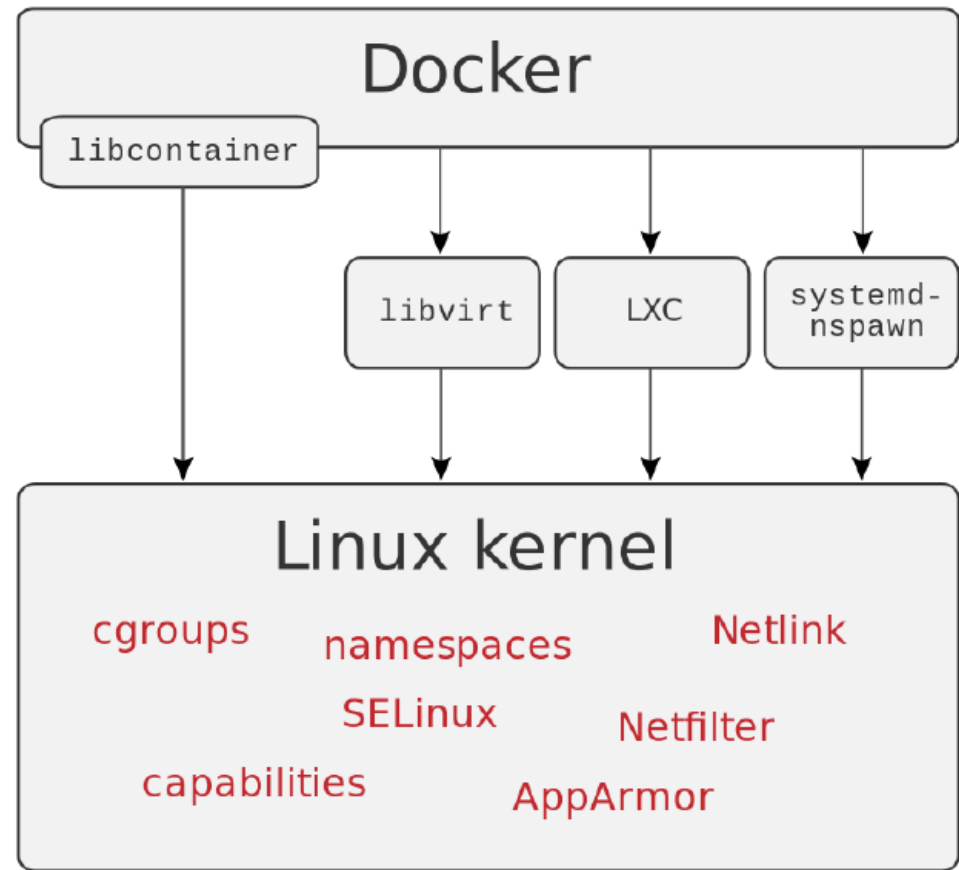


Docker internals

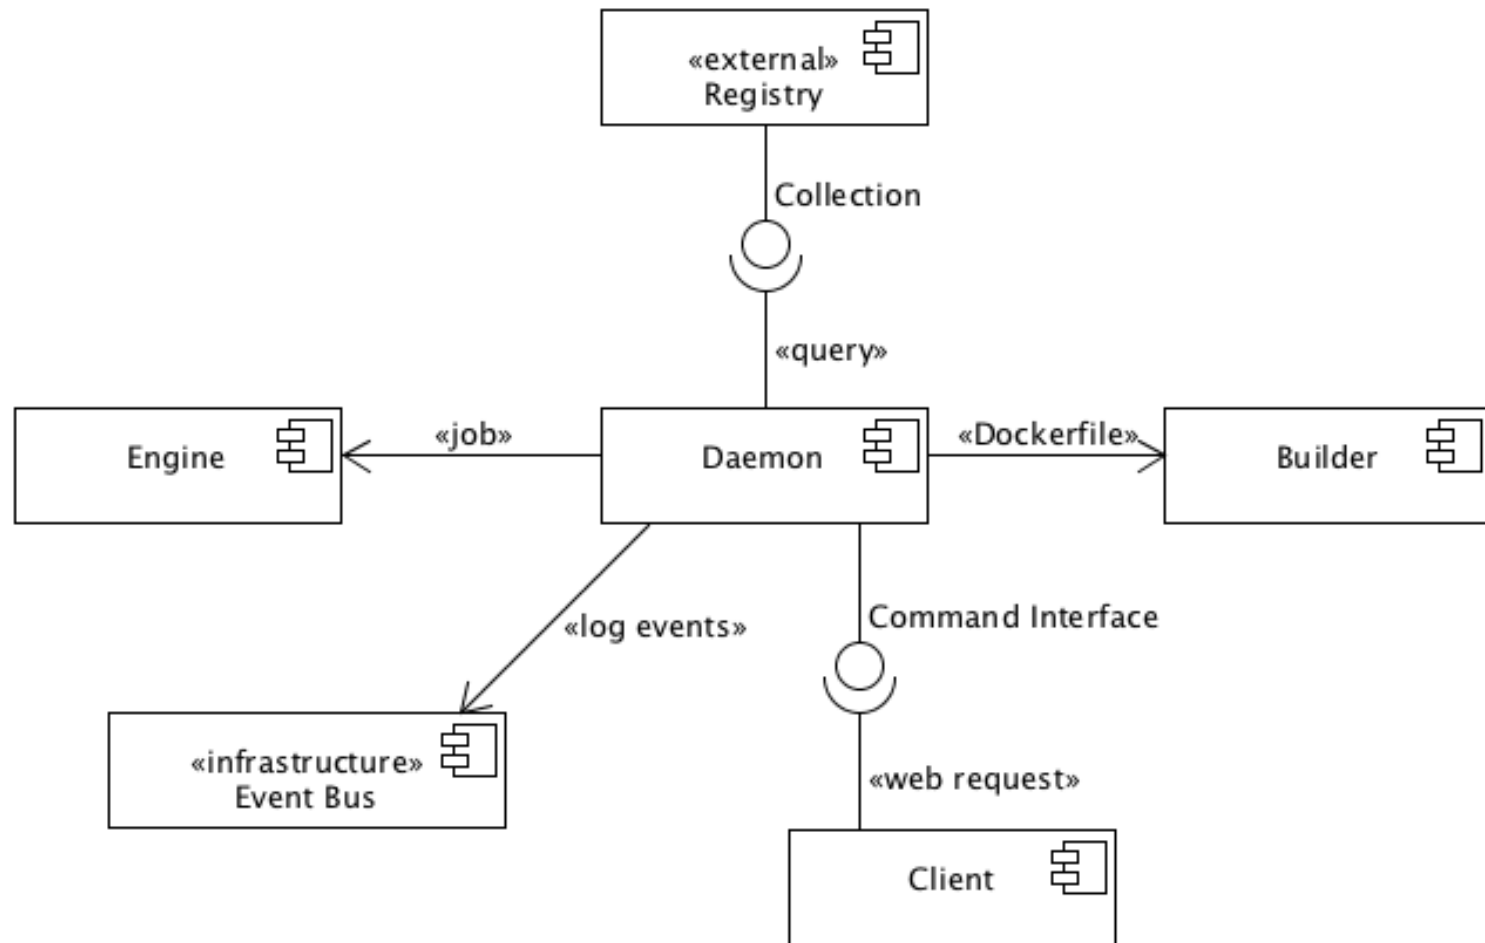
- ✿ Docker is written in Go language
- ✿ With respect to other OS-level virtualization solutions, Docker is a higher-level platform that exploits Linux kernel mechanisms such as **cgroups** and **namespaces**
 - ✦ First versions based on Linux Containers (LXC)
 - ✦ Then based on its own *libcontainer* runtime that uses Linux kernel *namespaces* and *cgroups* directly
- ✿ Docker adds to LXC
 - ✦ Portable deployment across machines
 - ✦ Versioning, i.e., git-like capabilities
 - ✦ Component reuse
 - ✦ Shared libraries, see Docker Hub hub.docker.com

Docker internals

- ✦ *libcontainer* (now included in *opencontainers/runc*): cross-system abstraction layer aimed to support a wide range of isolation technologies



Docker component diagram

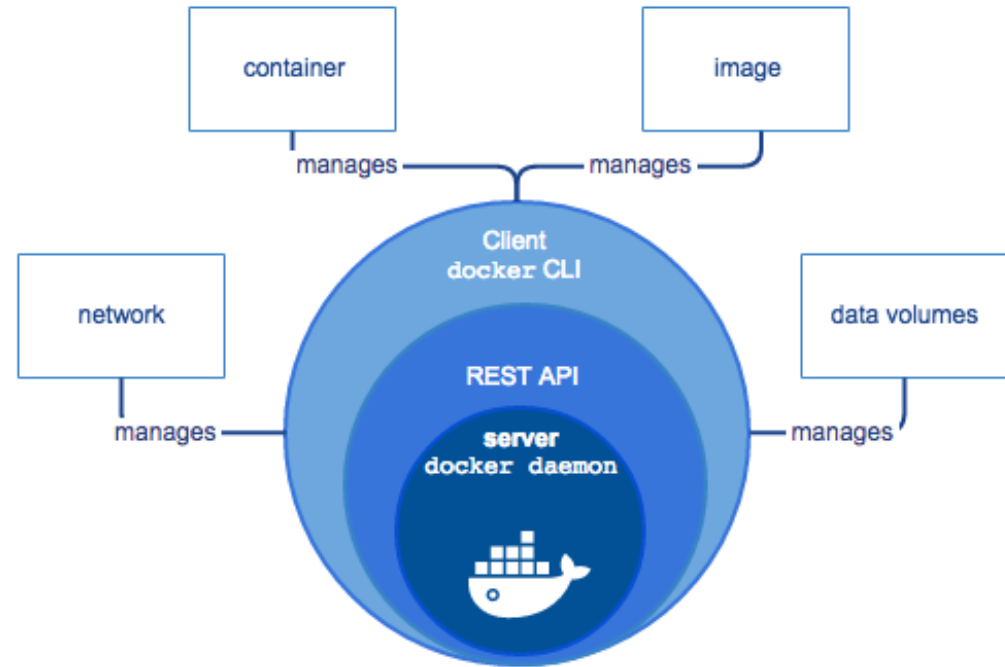


Docker

Docker engine

❖ ***Docker Engine***: client-server application composed by:

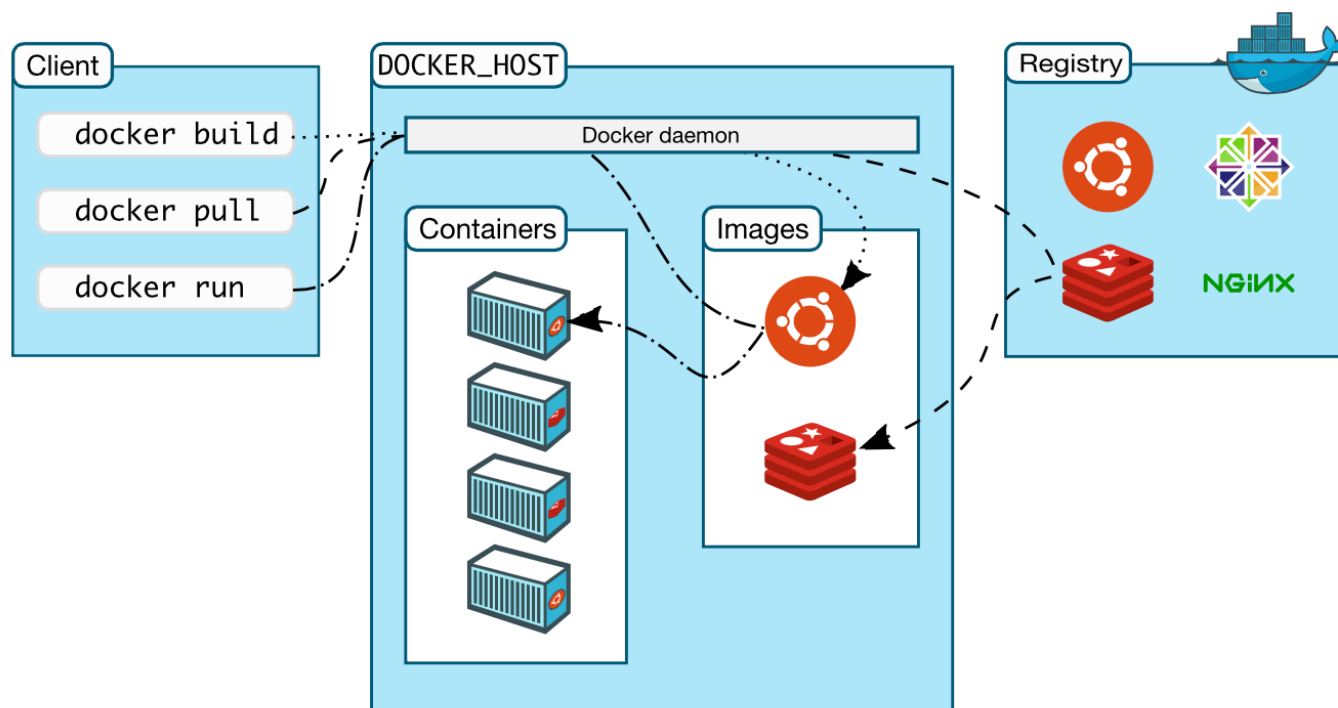
- ❖ A server, called docker daemon
- ❖ A REST API which specifies interfaces that programs can use to control and interact with the daemon
- ❖ A command line interface (CLI) client



See <https://docs.docker.com/engine/docker-overview/>

Docker architecture

- ❖ Docker uses a client-server architecture
 - ❖ The Docker *client* talks to the Docker *daemon*, which **builds**, **runs**, and **distributes** Docker containers
 - ❖ Client and daemon communicate via sockets or REST API



Docker



Docker image

- Read-only template used to create a Docker container
- The **Build** component of Docker
 - Enables the distribution of apps with their runtime environment
 - Incorporates all the dependencies and configuration necessary to apps to run, eliminating the need to install packages and troubleshoot
 - Target machine must be Docker-enabled
- Docker can build images automatically by reading instructions from a **Dockerfile**
 - A text file with simple, well-defined syntax
- Images can be pulled and pushed towards a public/private registry
- Image name: `[registry/][user/]name[:tag]`
 - Default for tag is **latest**



Docker image: Dockerfile

- Image can be created from a Dockerfile and a context
 - Dockerfile: instructions to assemble the image
 - Context: set of files (e.g., application, libraries)
- Often, an image is based on another image (e.g., ubuntu)
- Dockerfile syntax
 - # Comment
 - INSTRUCTION arguments
- Instructions in a Dockerfile run in order
- Some instructions
 - FROM**: to specify parent image (mandatory)
 - RUN**: to execute any command in a new layer on top of current image and commit results
 - ENV**: to set environment variables
 - EXPOSE**: container listens on specified network ports at runtime
 - CMD**: to provide defaults for executing container



Docker image: Dockerfile

- Example of Dockerfile to build the image of a container that will run a Python app (<https://docs.docker.com/get-started/>)

```
# Use an official Python runtime as a parent image
FROM python:2.7-slim

# Set the working directory to /app
WORKDIR /app

# Copy the current directory contents into the container at /app
ADD . /app

# Install any needed packages specified in requirements.txt
RUN pip install --trusted-host pypi.python.org -r requirements.txt

# Make port 80 available to the world outside this container
EXPOSE 80

# Define environment variable
ENV NAME World

# Run app.py when the container launches
CMD ["python", "app.py"]
```

Docker



Docker image: build

- ✿ Build image from Dockerfile


```
$ docker build [OPTIONS] PATH | URL | -
```

- ✿ E.g., to build the image for Python app (see Dockerfile in previous slide)

```
$ docker build -t friendlyhello .
```

Docker image: layers

- Each image consists of a *series of layers*
- Docker uses *union file systems* to combine these layers into a single unified view
 - Layers are stacked on top of each other to form a base for a container's root file system
 - Based on *copy-on-write* (COW) principle

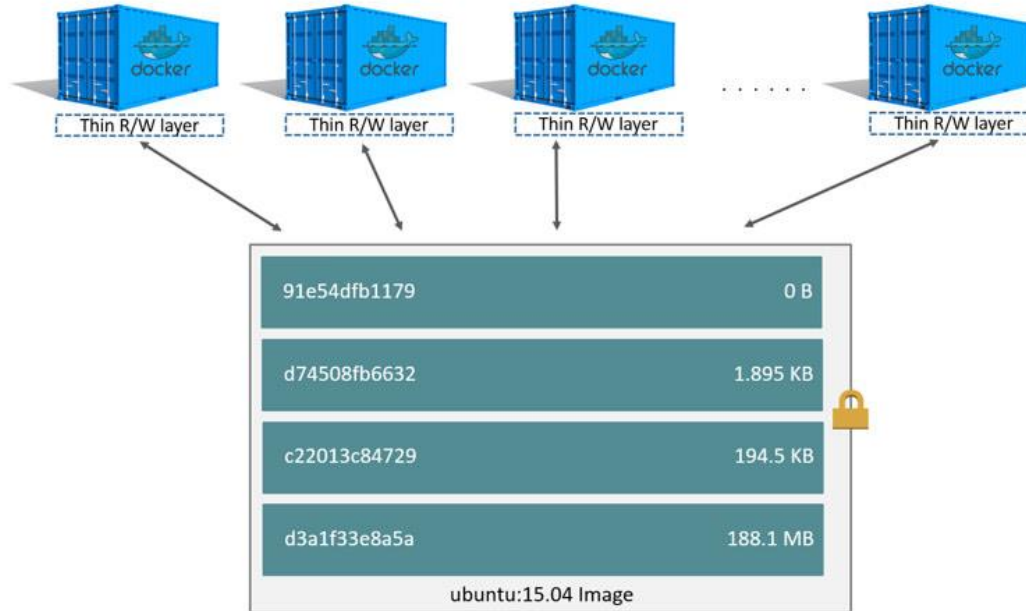


Docker image: layers

Layering pros

- Enable layer sharing and reuse, installing common layers only once and saving bandwidth and storage space
- Manage dependencies and separate concerns
- Facilitate software specializations

See <https://docs.docker.com/storage/storagedriver/>



Docker



Docker image: layers and Dockerfile

- Each layer represents an instruction in the image's Dockerfile
- Each layer except the very last one is read-only
- To inspect an image, including image layers
`$ docker inspect imageid`




Docker image: storage

- ✿ Containers should be **stateless**. Ideally:
 - ✦ Very little data is written to container's writable layer
 - ✦ Data should be written on **Docker volumes**
- ✿ Nevertheless: some workloads require to write data to the container's writable layer
- ✿ The **storage driver** controls how *images* and *containers* are stored and managed on the Docker host
- ✿ Multiple choices for the storage driver
 - ✦ Including AuFS and Overlay2 (at file level), Device Mapper, btrfs and zfs (at block level)
 - ✦ Storage driver's choice can affect the performance of containerized applications
 - ✦ See <https://docs.docker.com/storage/storagedriver/>

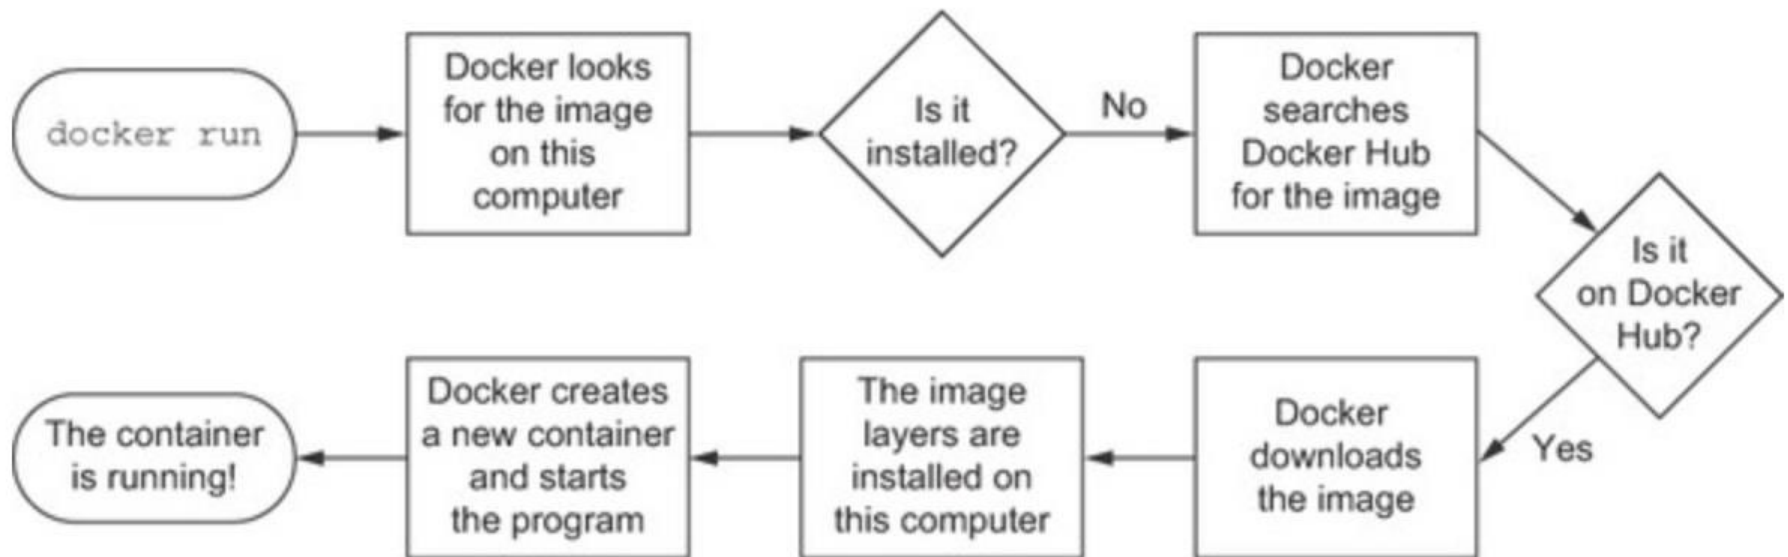
Docker container and registry

- ❖ **Docker container**: runnable instance of a Docker image
 - ❖ Run, start, stop, move, or delete a container using Docker API or CLI commands
 - ❖ The **Run** component of Docker
 - ❖ Docker containers are **stateless**: when a container is deleted, any data written not stored in a *data volume* is deleted along with the container
- ❖ **Docker registry**: stateless server-side application that stores and lets you distribute Docker images
 - ❖ Open library of images
 - ❖ The **Distribute** component of Docker
 - ❖ Docker-hosted registries: Docker Hub, Docker Store (open source and enterprise verified images)



Docker: run command

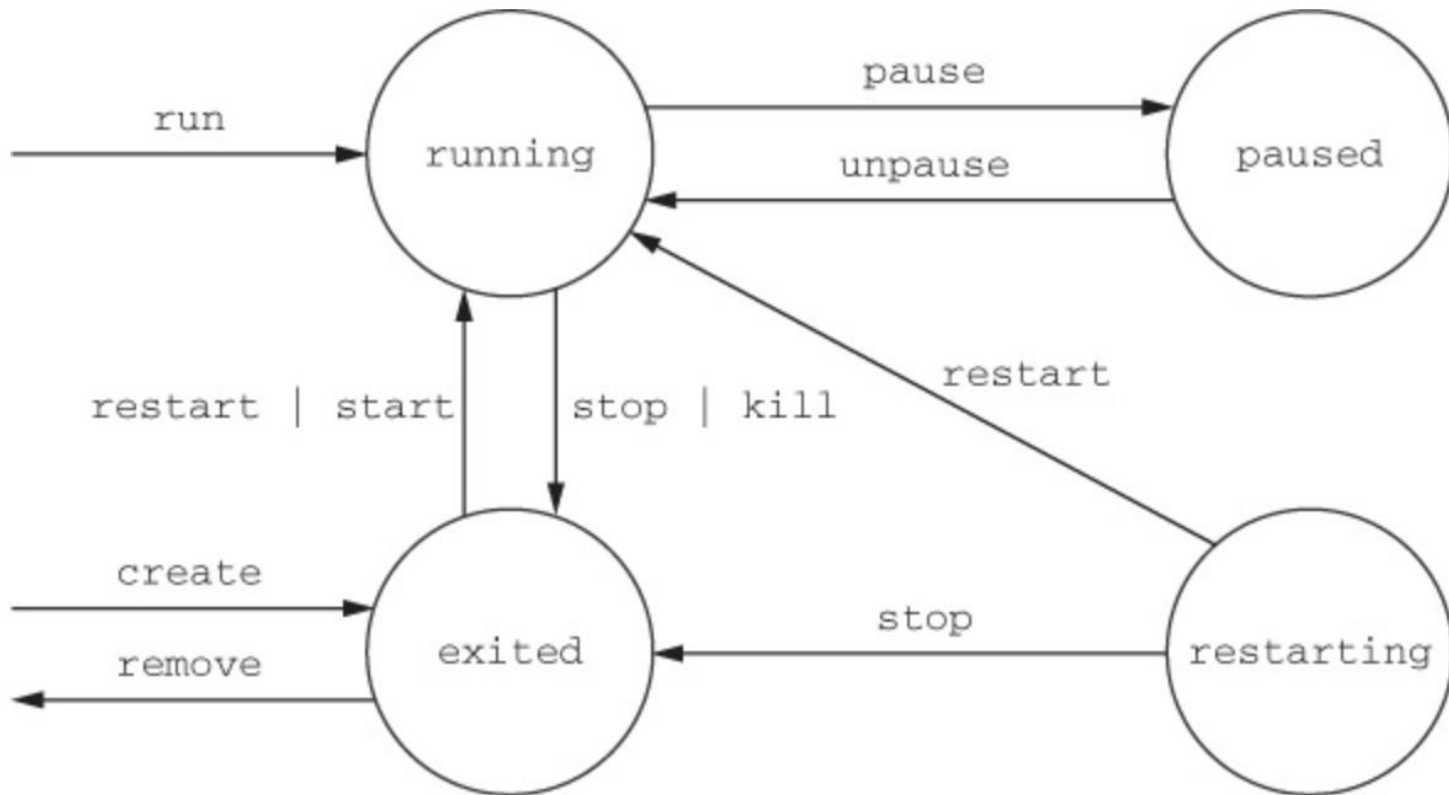
- When you run a container whose image is not yet installed but is available on Docker Hub



Docker in Action by Jeff Nickoloff

Docker containers

State transitions



Docker in Action by Jeff Nickoloff

Docker



Commands: Docker info

- Obtain system-wide info on Docker installation

```
$ docker info
```

- Including:

- How many images, containers and their status
- Storage driver
- Operating system, architecture, total memory
- Docker registry
- Docker Swarm status



Commands: image handling

- List images on host (i.e., local repository)
`$ docker images`
- List every image, including intermediate image layers:
`$ docker images -a`
- Options to list images by name and tag, to list image digests (sha256), to filter images, to format the output, e.g.,
`$ docker images --filter reference=ubuntu`
- Remove an image (can also use `imagename` instead of `imageid`)
`$ docker rmi imageid`



Command: run

```
$ docker run [OPTIONS] IMAGE [COMMAND] [ARGS]
```

Most common options

- name assign a name to the container
- d detached mode (in background)
- i interactive (keep STDIN open even if not attached)
- t allocate a pseudo-tty
- expose expose a range of ports inside the container
- p publish a container's port or a range of ports to the host
- v bind and mount a volume
- e set environment variables
- link add link to other containers

The "Hello World" container

```
$ docker run alpine /bin/echo 'Hello world'
```

- alpine: lightweight Linux distro with reduced image size



Commands: containers management

✚ List containers

- ✚ Only running containers: `$ docker ps`
- ✚ Alternatively, `$ docker container ls`
- ✚ All containers (even stopped or killed containers):

`$ docker ps -a`

✚ Container lifecycle

- ✚ Stop running container
`$ docker stop containerid`
- ✚ Start stopped container
`$ docker start containerid`
- ✚ Kill running container
`$ docker kill containerid`
- ✚ Remove container (need to stop it before attempting removal)
`$ docker rm containerid`

Can also use `containername` instead of `containerid`



Commands: containers management

Inspect a container

- Most detailed view of the environment in which a container was launched

```
$ docker inspect containerid
```

Copy files from and to docker container

```
$ docker cp containerid:path localpath
```

```
$ docker cp localpath containerid:path
```



Examples of using Docker

- Run a nginx Web server inside a container
 - Also bind the container to a specific port
- Send HTTP request through Web browser
 - First retrieve the hostname of the host machine
- Send HTTP request through an interactive container using **Docker internal network**

```
$ docker run -i -t --link web:web --name web_test busybox
/ # wget -O - http://web:80/
/ # exit
```

--link: legacy flag to manually create links between the containers
wget: -O FILE Save to FILE ('-' for stdout)

- Instead of using --link, let us define a **bridge network**

```
$ docker network create my_net
$ docker run -d -p 80:80 --name web --net=my_net nginx
$ docker run -i -t --net=my-net --name web_test busybox
/ # ...
```




Examples of using Docker

- Send HTTP request through an Alpine Linux container with curl installed and set as entrypoint

```
$ docker run --rm byrnedo/alpine-curl http://...
```

- Check container logs

```
$ docker logs containerid
```

Examples of using Docker

Running Apache web server with minimal index page

- Define container image with Dockerfile
 - Define image starting from Ubuntu, install and configure Apache
 - Incoming port set to 80 using EXPOSE instruction

```
FROM ubuntu

# Install dependencies
RUN apt-get update
RUN apt-get -y install apache2
# Install apache and write hello world message
RUN echo 'Hello World!' > /var/www/html/index.html
# Configure apache
RUN echo '. /etc/apache2/envvars' > /root/run_apache.sh
RUN echo 'mkdir -p /var/run/apache2' >> /root/run_apache.sh
RUN echo 'mkdir -p /var/lock/apache2' >> /root/run_apache.sh
RUN echo '/usr/sbin/apache2 -D FOREGROUND' >> /root/run_apache.sh
RUN chmod 755 /root/run_apache.sh

EXPOSE 80

CMD /root/run_apache.sh
```



Examples of using Docker

- Build container image from Dockerfile

```
$ docker build -t hello-apache .
```

- Run container and bind

```
$ docker run -d -p 80:80 hello-apache
```

- Option `-p`: publish container port (80) to host port (80)



Examples of using Docker

✿ Stop and remove a container

```
$ docker ps
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
2b2b39d05fe0	hello_world	"/bin/sh -c /root/ru..."	3 seconds ago	Up 2 seconds	0.0.0.0:80->80/tcp	pedantic_austin

```
$ docker stop containerid
```

```
$ docker ps -a
```

```
$ docker rm containerid
```

✿ Stop all containers

```
$ for i in $(docker ps -q); do docker stop $i; done
```



Examples of using Docker

Running a Python web app in Docker

- ❏ See <https://docs.docker.com/get-started/part2/>
- ❏ Define container image with Dockerfile
 - Define image with Python runtime and Python app
 - Incoming port set to 80 using EXPOSE instruction
- ❏ Write app code in Python using Flask and Redis packages and create file requirements.txt to specify those packages needed by app
 - Redis: in-memory data store, used to keep the counter of the number of visits to web app
- ❏ Build container image

```
$ docker build -t friendlyhello .
```
- ❏ Run container and bind

```
$ docker run -d -p 4000:80 friendlyhello
```

 - Option `-p`: publish container port (80) to host port (4000)



Multi-container Docker applications

🔗 How to run multi-container Docker apps?

1. Docker Compose

- ❏ Deployment only on single host

2. Docker Swarm


- ❏ Native orchestration tool for Docker
- ❏ Deployment on multiple hosts

3. Kubernetes

- ❏ Deployment on multiple hosts



Docker Compose



- ✿ To coordinate the execution of multiple containers, we can use **Docker Compose**
 - ✦ See <https://docs.docker.com/compose/>
- ✿ Docker Compose
 - ✦ Not bundled within Docker installation (on Linux)
 - ✦ <https://docs.docker.com/compose/install/>
- ✿ Allows to easily express the containers to be instantiated at once, and the relations among them
- ✿ Runs the composition on a single machine (i.e., single Docker engine)
- ✿ Use **Docker Swarm** if you need to deploy containers on multiple nodes

Docker Compose

- Specify how to compose containers in a easy-to-read file named `docker-compose.yml`

- To start Docker composition (background -d):

```
$ docker-compose up -d
```

- To stop Docker composition:

```
$ docker-compose down
```

- By default, Docker Compose looks for `docker-compose.yml` in current working directory
 - Change file using -f flag

Docker Compose

- ❁ Different versions of the Docker compose file format
 - ❁ Latest: version 3 is supported from Docker Compose 1.13

```
version: '3'

services:
  storm-nimbus:
    image: storm
    container_name: nimbus
    command: storm nimbus
    depends_on:
      - zookeeper
    links:
      - zookeeper
    ports:
      - "6627:6627"
```



```
zookeeper:
  image: zookeeper
  container_name: zookeeper
  ports:
    - "2181:2181"

worker1:
  image: storm
  command: storm supervisor
  depends_on:
    - storm-nimbus
    - zookeeper
  links:
    - storm-nimbus
    - zookeeper
```

Docker Compose file format:

<https://docs.docker.com/compose/compose-file/>




Docker Compose: example

- ✿ Simple Python web app running on Docker Compose
 - ✦ Two containers: Python web app and Redis
 - ✦ Use Flask framework and maintain a hit counter in Redis
 - ✦ See <https://docs.docker.com/compose/gettingstarted/>
- ✿ Steps:
 - ✦ Write Python app
 - ✦ Define Python container image with Dockerfile
 - ✦ Define services in [docker-compose.yml](#) file
 - Two services: web (image defined by Dockerfile) and redis (image pulled from Docker Hub)
 - ✦ Build and run your app with Compose

```
$ docker-compose up -d
```
 - ✦ Send HTTP requests using curl ([now counter is increased](#))
 - ✦ Stop Compose

```
$ docker-compose down
```

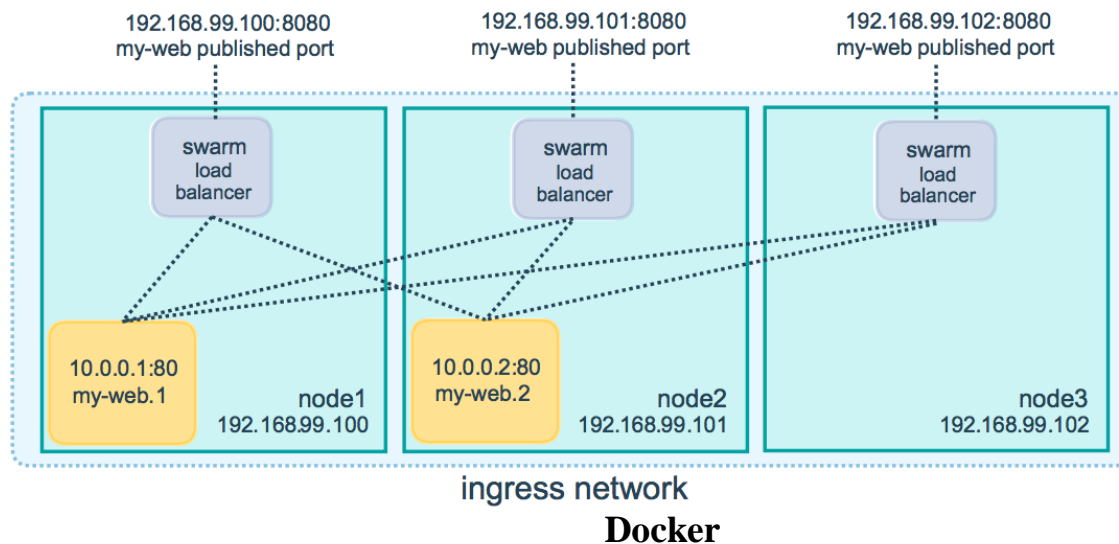
Docker Swarm



- ❁ Docker includes **swarm mode** for natively managing a cluster of Docker Engines, called *swarm*
 - ❁ See <https://docs.docker.com/engine/swarm/>
- ❁ **Tasks:** containers running in a **service**
- ❁ Basic features of swarm mode:
 - ❁ **Scaling:** number of tasks for each service
 - ❁ **State reconciliation:** Swarm monitors cluster state and reconciles any differences w.r.t. desired state (e.g., replace containers after host failure)
 - ❁ **Multi-host networking:** to specify an overlay network among services
 - ❁ **Load balancing:** allows to expose the ports for services to an external load balancer. Internally, the swarm lets you specify how to distribute containers among nodes

Docker Swarm

- ✿ A swarm consists of multiple Docker hosts which run in swarm mode
- ✿ Node: instance of Docker engine
 - ✦ **Manager node** dispatches tasks to worker nodes
 - ✦ **Worker nodes** receive and execute tasks
- ✿ Load balancing
 - ✦ Swarm manager can automatically assign the service a (configurable) PublishedPort
 - ✦ External components can access the service on PublishedPort. All nodes in the swarm route ingress connections to a running task





Docker Swarm: Swarm cluster

✿ Create a swarm: manager node

```
$ docker swarm init --advertise-addr <MANAGER-IP>  
Swarm initialized: current node (<nodeid>) is now a manager.  
To add a worker to this swarm, run the following command:
```

```
docker swarm join --token <token> <manager-ip>:port
```

✿ Create a swarm: worker node

```
$ docker swarm join --token <token> <manager-ip>:port
```

✿ Inspect status

```
$ docker info
```

```
$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
<nodeid> *	controller	Ready	Active	Leader
<nodeid>	storage	Ready	Active	



Docker Swarm: Swarm cluster

✦ Leave the swarm

```
$ docker swarm leave
```

- ✦ If the node is a manager node, warning about maintaining the quorum (to override warning, --force flag)

✦ After a node leaves the swarm, you can run docker node rm command on a manager node to remove the node from the node list

```
$ docker node rm node-id
```



Docker Swarm: manage services

- Deploy a service to the swarm (from manager node)

```
$ docker service create -d --replicas 1 \  
  --name helloworld alpine ping docker.com
```

- List running services

```
$ docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
<serviceid>	helloworld	replicated	1/1	alpine:latest	



Docker Swarm: manage services

Inspect service

```
$ docker service inspect --pretty <SERVICE-ID>
$ docker service ps <SERVICE-ID>
```

ID	NAME	IMAGE	NODE	DESIRED	ST	CURRENT	ST	ERROR	PORTS
<cont.id1>	helloworld.1	alpine:latest	controller	Running		Running	...		
<cont.id2>	helloworld.2	alpine:latest	storage	Running		Running	...		

Inspect container

```
$ docker ps <cont.id1>
```

Manager node

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	... NAMES
<cont.id1>	alpine:latest	"ping docker.com"	2 min ago	Up 2 min	helloworld.1.iuk1sj...

Worker node

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	... NAMES
<cont.id2>	alpine:latest	"ping docker.com"	2 min ago	Up 2 min	helloworld.2.skfos4...

Docker Swarm: manage services

Scale service

```
$ docker service scale <SERVICE-ID>=<NUMBER-OF-TASKS>
```

- Swarm manager will automatically enact the updates

Apply rolling updates to a service

```
$ docker service update --limit-cpu 2 redis  
$ docker service update --replicas 2 helloworld
```

Roll back an update

```
$ docker service rollback [OPTIONS] <SERVICE-ID>
```

Remove a service

```
$ docker service rm <SERVICE-ID>
```

References

- ✿ Docker Documentation

- ✿ <https://docs.docker.com/>
- ✿ <https://docs.docker.com/get-started/>

- ✿ Docker for beginners

- ✿ <https://docker-curriculum.com/>

- ✿ Ian Miell and Aidan Hobson Sayers, *Docker in Practice*, 2nd Edition, Manning Publications, 2019