



Internet of Things and Services

Service-oriented architectures

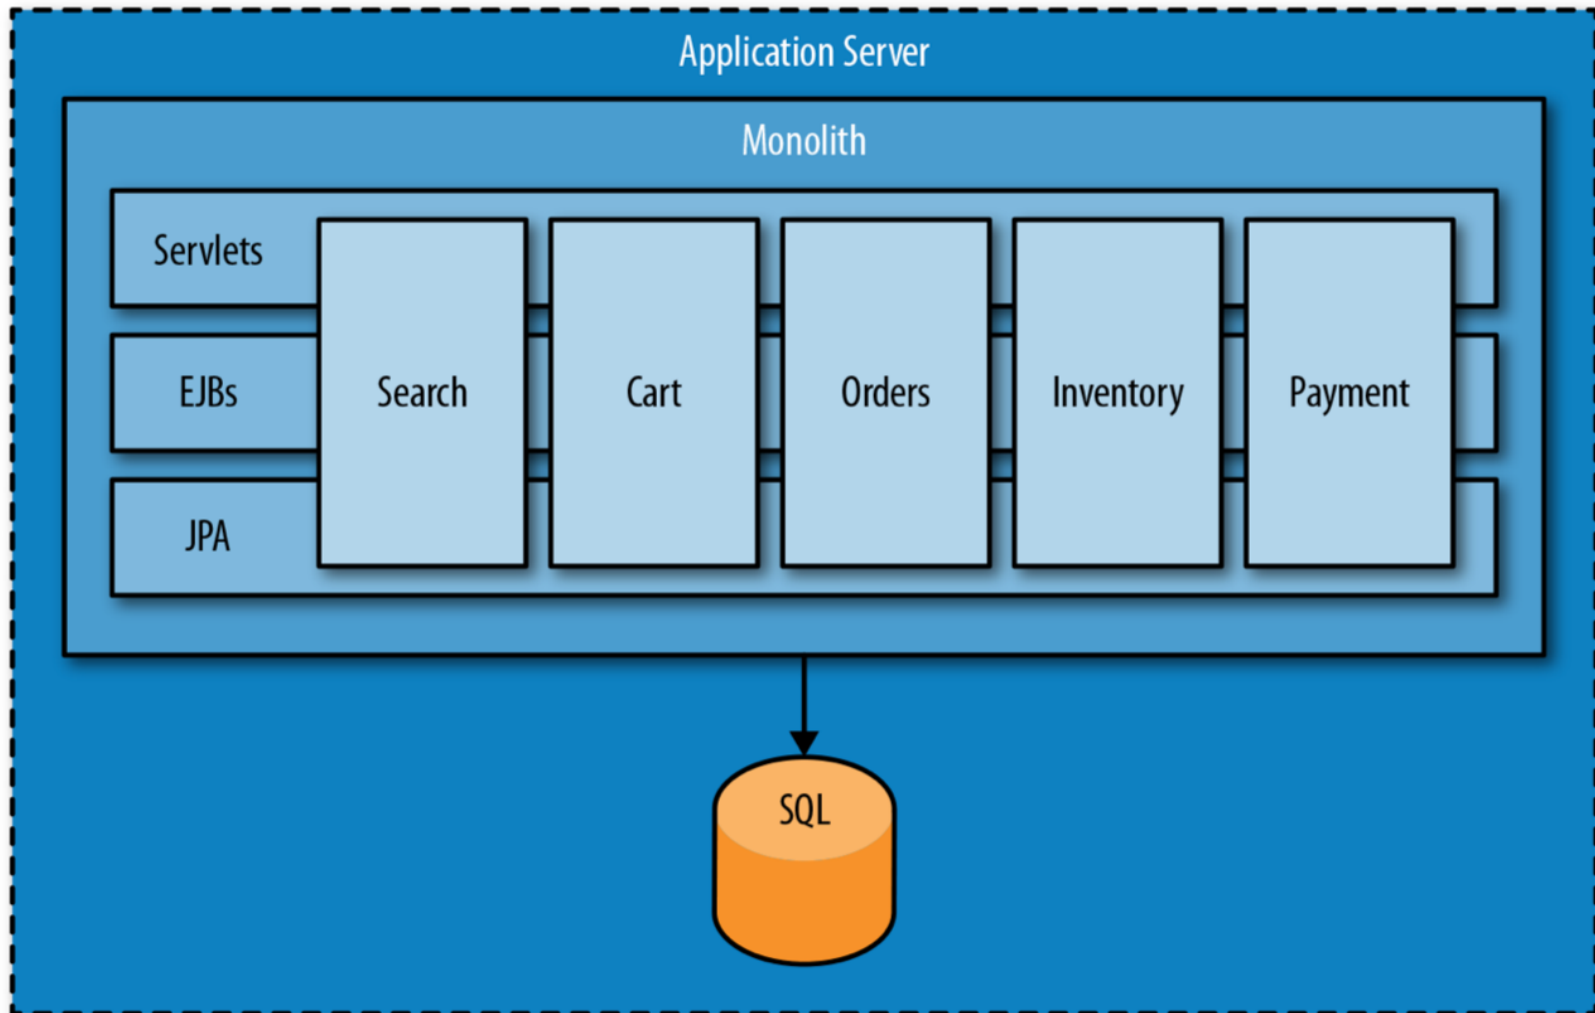
Reactive microservices and serverless computing

Department of Computer Science
Faculty of Electronic Engineering, University of Nis

Internet of Things and Services
Computing and informatics

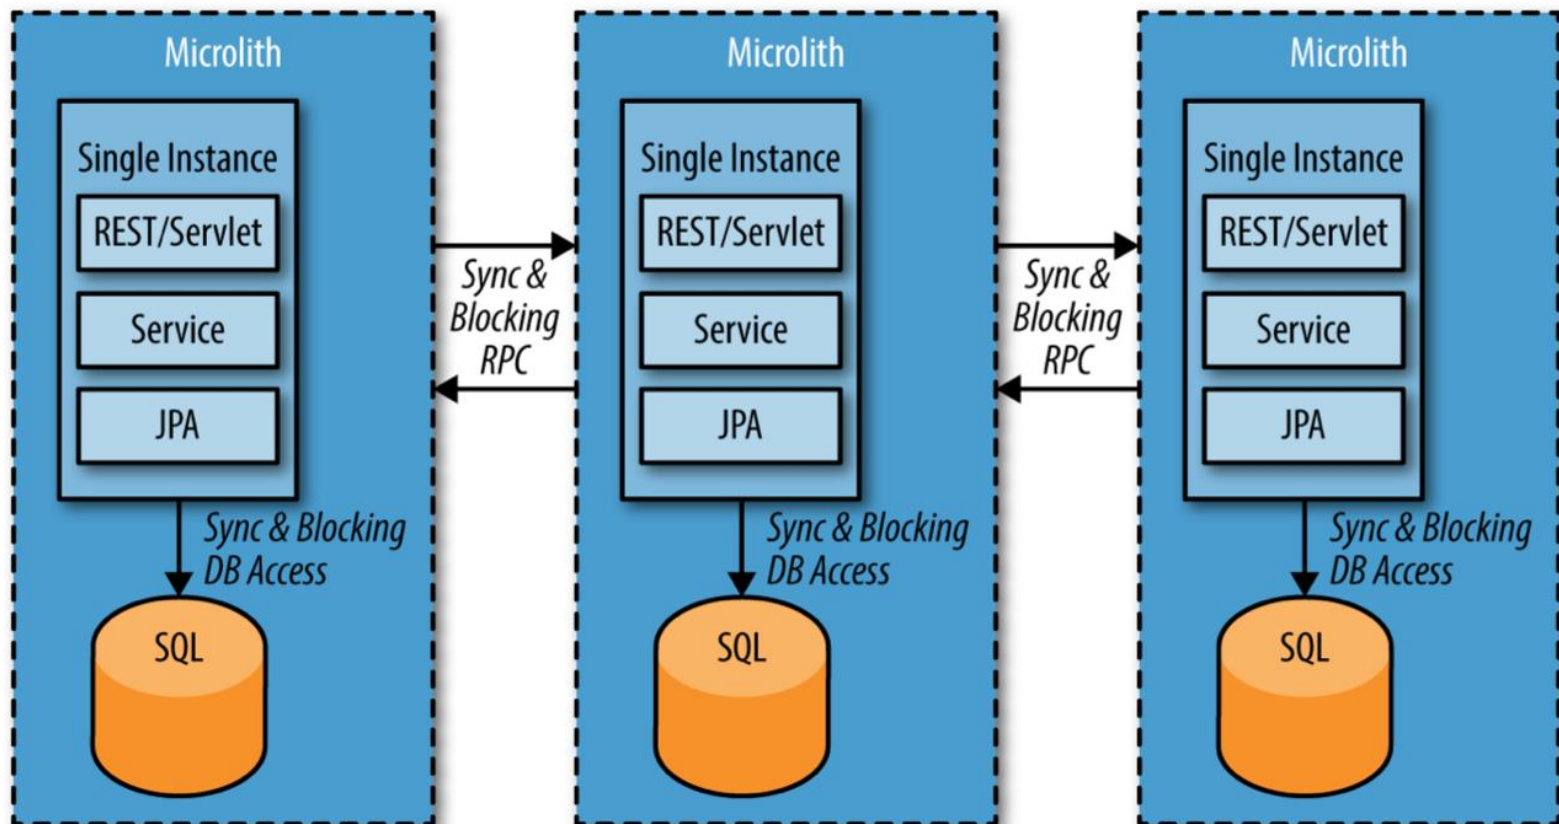
Prof. dr Dragan Stojanović

Monolith architecture



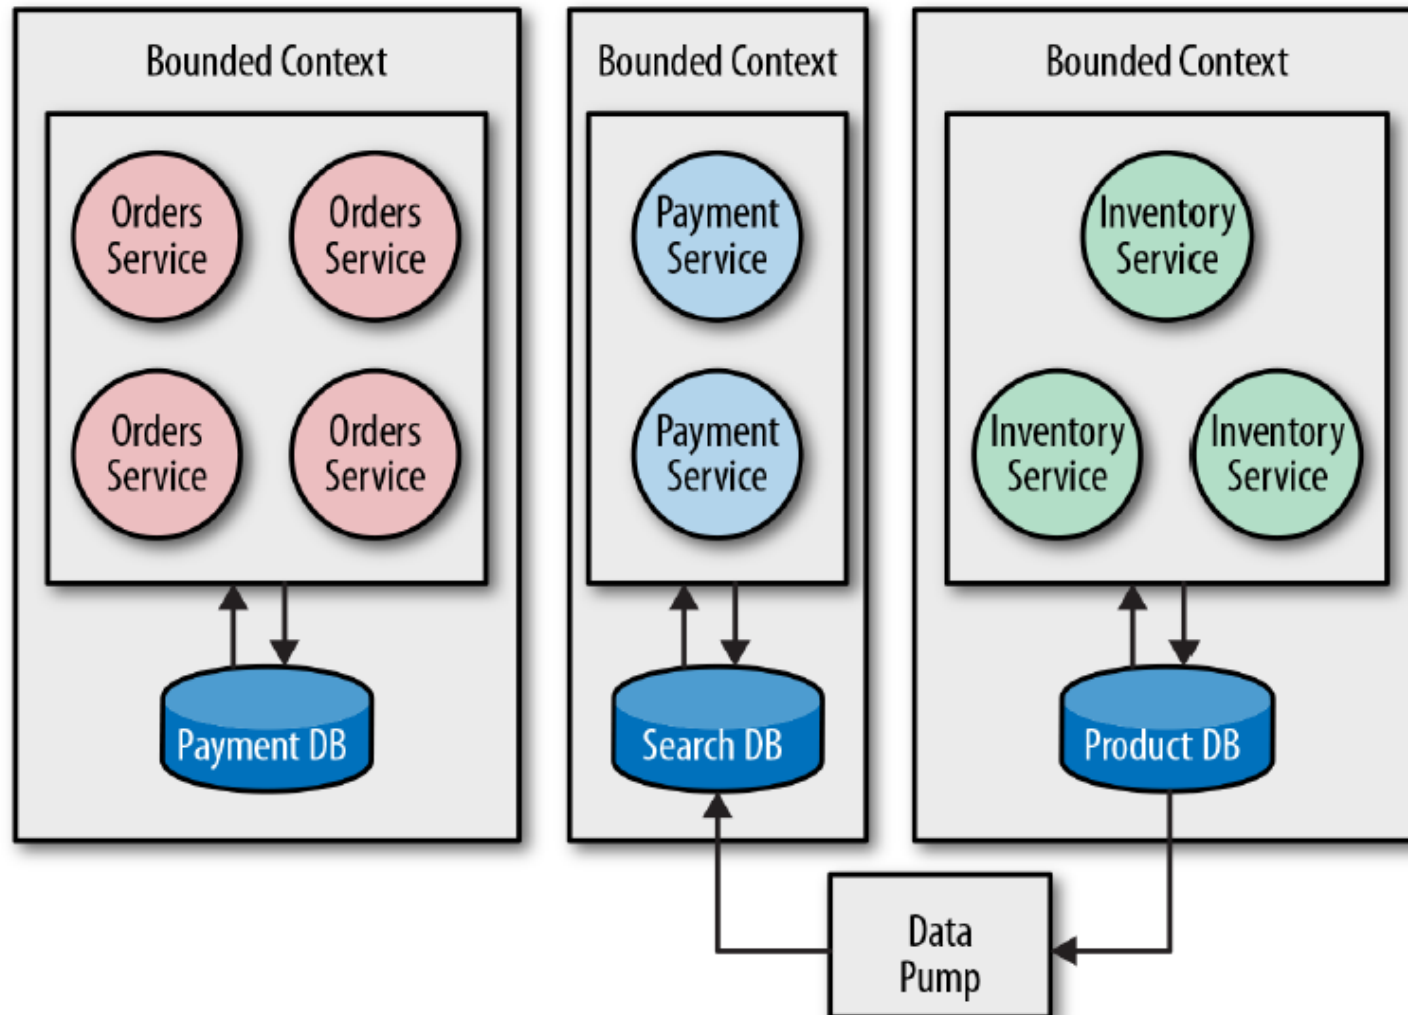
Microlith architecture

- Single instance microservices communicating over blocking protocols – **Not resilient, not scalable**



Reactive microservices and serverless computing

Microservice architecture



Reactive microservices and serverless computing

Internet of Things and Services



Essential traits of a microservice (1)

✿ Isolation

- ✿ A prerequisite for resilience and elasticity and requires asynchronous communication boundaries between services to decouple them in: Time (allowing concurrency) and Space (allowing distribution and mobility—the ability to move services around)
- ✿ Isolation also makes it easier to scale each service, as well as allowing them to be monitored, debugged and tested independently

✿ Autonomicity

- ✿ Only when services are isolated can they be fully autonomous and make decisions independently, act independently, and cooperate and coordinate with others to solve problems.

✿ Single responsibility

- ✿ Do One Thing, and Do It Well



Essential traits of a microservice (2)

❁ Exclusive state (Own Your State, Exclusively)

- ❁ Responsibility for their own state and the persistence thereof.
- ❁ When communicating with another Microservice, across Bounded Contexts, each service responds to a request at its own will, with immutable data (facts) derived from its current state, and never exposes its mutable state directly.
- ❁ If the Command to the service triggers a state change in the service then we can capture the state change as a new fact in an Event to be stored in the Event Log using Event Sourcing.

❁ Asynchronous Message-Passing

- ❁ Asynchronous boundary between services is necessary in order to decouple them, and their communication flow, in *time* and *space*

❁ Mobility


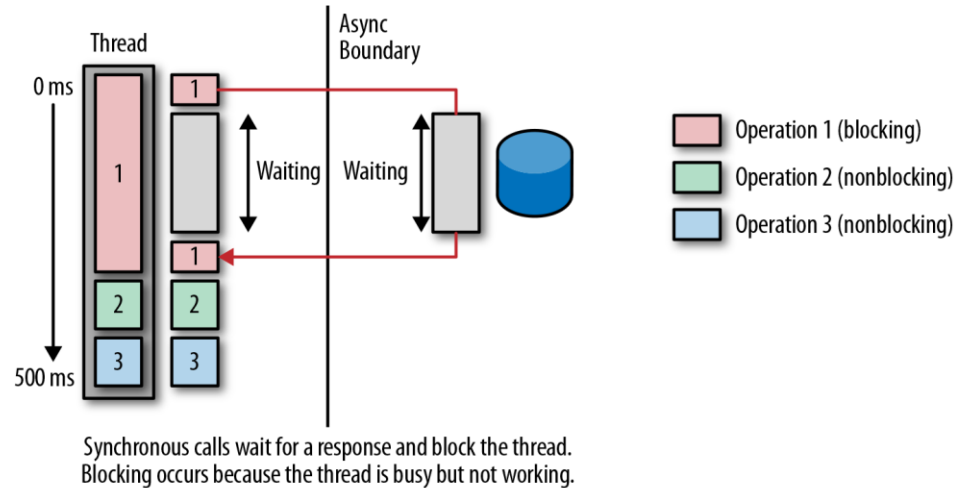
- ❁ Decoupling in space, *Location Transparency*, the ability to, at runtime, dynamically scale the Microservice, either on multiple cores or on multiple nodes, without changing the code

Reactive Programming

- **Reactive Programming** is essential to the design of microservices, allowing us to build highly efficient, responsive, and scalable services.

- **Techniques for Reactive Programming include:**

- ❏ Asynchronous execution and I/O
- ❏ Back-pressured streaming,
- ❏ Circuit breakers.



Reactive Programming


- ✿ Many of the ideas behind reactive are not new; plenty of them were described and implemented years ago, for example, Erlang's actor-based programming model (early 1980s), and Akka (JVM).
- ✿ The need for concurrent and distributed applications is growing stronger.
- ✿ A number of movements have contributed to the current rise of reactive programming, most notably:
 - ✦ **IoT and mobile systems and applications**
 - The server side has to handle millions of connected devices concurrently, a task best handled by asynchronous processing, due to its lightweight ability to represent resources such as “the device,” or whatever it might be.
 - ✦ **Cloud and containerization**
 - While we've had cloud-based infrastructure for a number of years now, the rise of lightweight virtualization and containers, together with container-focused schedulers and PaaS solutions, has given us the freedom and speed to deploy much faster and with a finer-grained scope.

Reactive Programming

- ⊗ ReactiveX extensions (API) for many programming languages - <http://reactivex.io/>
 - ⊗ Java (RxJava), C# (Rx.NET), C++ (RxCpp), Python (RxPY), JavaScript (RxJS), Go (RxGo),...
- ⊗ Reactive frameworks on top of the JDK (JVM based) and implement the Reactive Streams specification
 - ⊗ Reactor, Vert.x, Akka, Ratpack,...
- ⊗ Programming languages that have native reactive capabilities
 - ⊗ Java 9 and Spring Framework 5, and Clojure, Scala, Go,...
- ⊗ JavaScript libraries
 - ⊗ Angular.js, React.js, Ractive.js, and Node.js that can be used to build front-end reactive applications.

Toward Reactive Microsystems (1)

- Embrace Reactive Programming
- Asynchronous and nonblocking I/O (Pub-Sub messaging)
- Applying backpressure
 - Backpressure is all about flow control, ensuring that a fast producer should not be able to overwhelm a slower consumer by being allowed to send it more data than it can handle



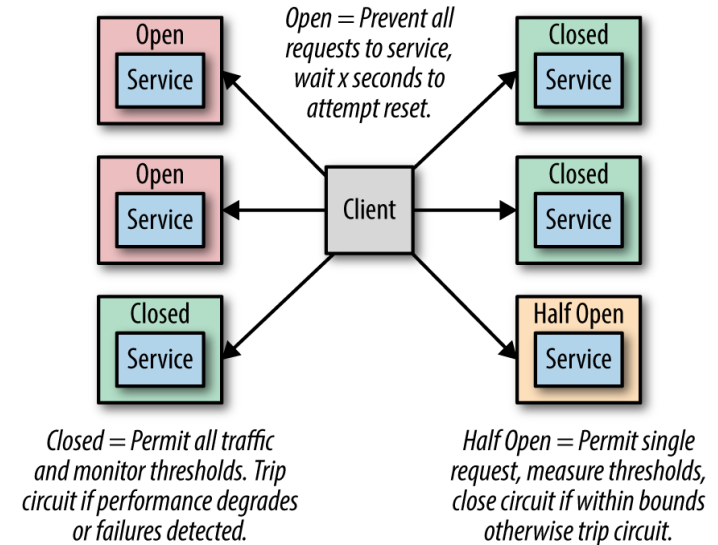
Toward Reactive Microsystems (2)




Circuit breaker

- A finite-state machine (FSM), which means that it has a finite set of states: *Closed*, *Open*, and *Half-Open*. The default state is *Closed*, which allows all requests to go through
- When a failure (or a specific number of failures) have been detected, the circuit breaker “trips” and moves to an *Open* state. In this state, it does not let any requests through, and instead fails fast to shield the component from the failed service.
- After a timeout has occurred, the service is back up again, so it attempts to “reset” itself and move to a *Half-Open* state

Circuit breakers act as a wrapper around external services



When an external service is too slow or begins to fail, a circuit breaker trips and prevents requests to that service...



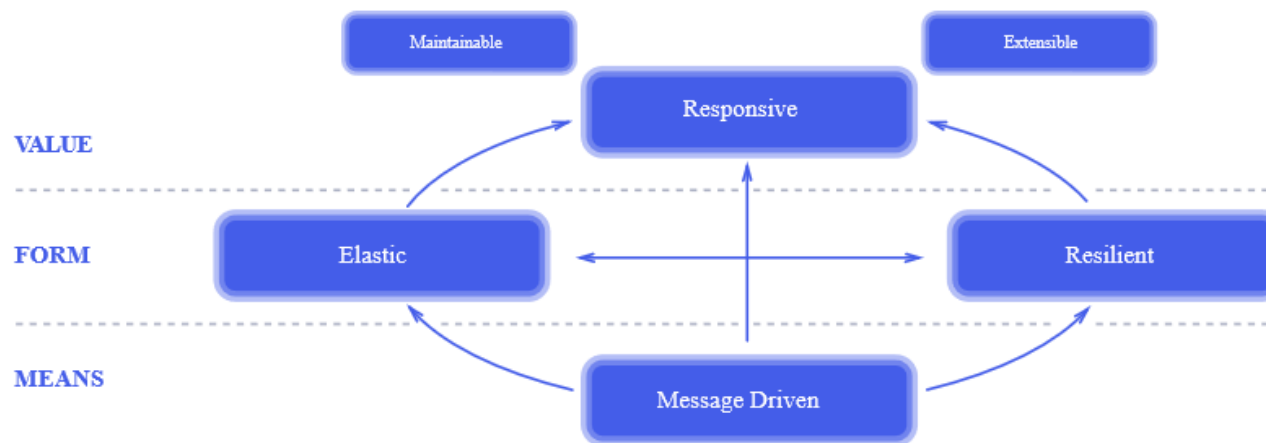
...until the service heals. This prevents cascading failures in a distributed system.

Reactive Systems (1)

- ✿ **Reactive Manifesto** - <https://www.reactivemanifesto.org/>
- ✿ **Reactive Systems are:**
 - ✦ **Responsive:** The **system** responds in a timely manner if at all possible. Responsive systems focus on providing rapid and consistent response times, establishing reliable upper bounds so they deliver a consistent quality of service.
 - ✦ **Resilient:** The system stays responsive in the face of **failure**. This applies not only to highly-available, mission-critical systems — any system that is not resilient will be unresponsive after a failure. Resilience is achieved by **replication**, containment, **isolation** and **delegation**. Failures are contained within each component, isolating components from each other and thereby ensuring that parts of the system can fail and recover without compromising the system as a whole.

Reactive Systems (2)

- ✿ **Elastic**: The system stays responsive under varying workload. Reactive Systems can react to changes in the input rate by increasing or decreasing the resources allocated to service these inputs.
- ✿ **Message Driven**: Reactive Systems rely on **asynchronous message-passing** to establish a boundary between components that ensures loose coupling, isolation and location transparency. Employing explicit message-passing enables load management, elasticity, and flow control by shaping and monitoring the message queues in the system and applying **back-pressure** when necessary.



Reactive microservices and serverless computing

Internet of Things and Services



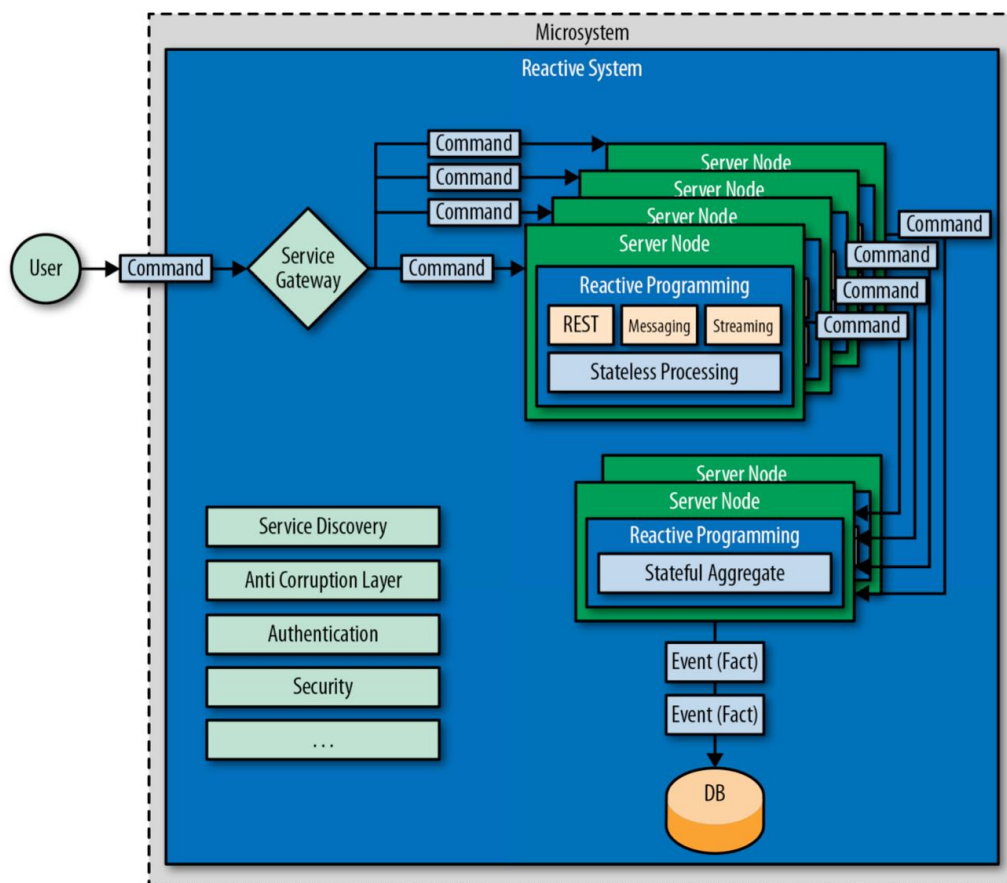
Reactive programming vs. Reactive systems

- ✿ **Reactive Programming** is a great technique for making individual components performant and efficient through **asynchronous** and **non-blocking** execution, most often together with a mechanism for **backpressure**.
 - ✦ It has a *local focus* and is *event-driven*, publishing facts to 0–N anonymous subscribers.
- ✿ **Reactive Systems** takes a holistic view on system design, focusing on keeping *distributed systems* responsive by making them *resilient and elastic*.
 - ✦ It is *message-driven*, based upon asynchronous message passing, which makes distributed communication to addressable recipients first class—allowing for elasticity, location transparency, isolation, supervision, and self-healing.

<https://www.oreilly.com/radar/reactive-programming-vs-reactive-systems/>

Microservices as Reactive Systems

- Each microservice needs to be designed as a component of a distributed system - a **Microsystem**



Reactive microservices and serverless computing

Internet of Things and Services

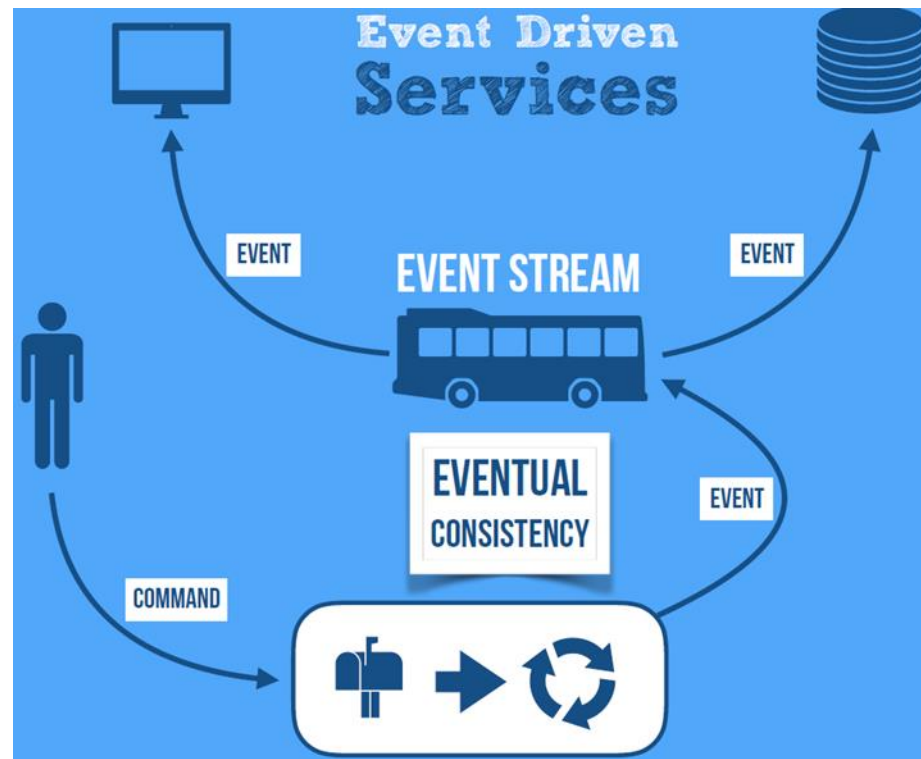


Event Driven Architecture

- ✿ **Event-Driven Architecture** (EDA) is a design paradigm in which a software component executes in response to receiving one or more event notifications
- ✿ **Events** represent **Facts** of information
 - ✦ Facts are immutable
 - ✦ Facts accrue - Knowledge can only grow
- ✿ Events/Facts can be **Disregarded/Ignored**
- ✿ Events/Facts can not be **Retracted** (once accepted)
- ✿ Events/Facts can not be **Deleted** (once accepted)
- ✿ Events/Facts (new) can **Invalidate** existing Facts

Event Driven Services

1. Receive and react (or not) to facts, that are coming its way
2. Publish new facts (as events) to the rest of the world
3. Invert the control flow to minimize coupling and increase autonomy

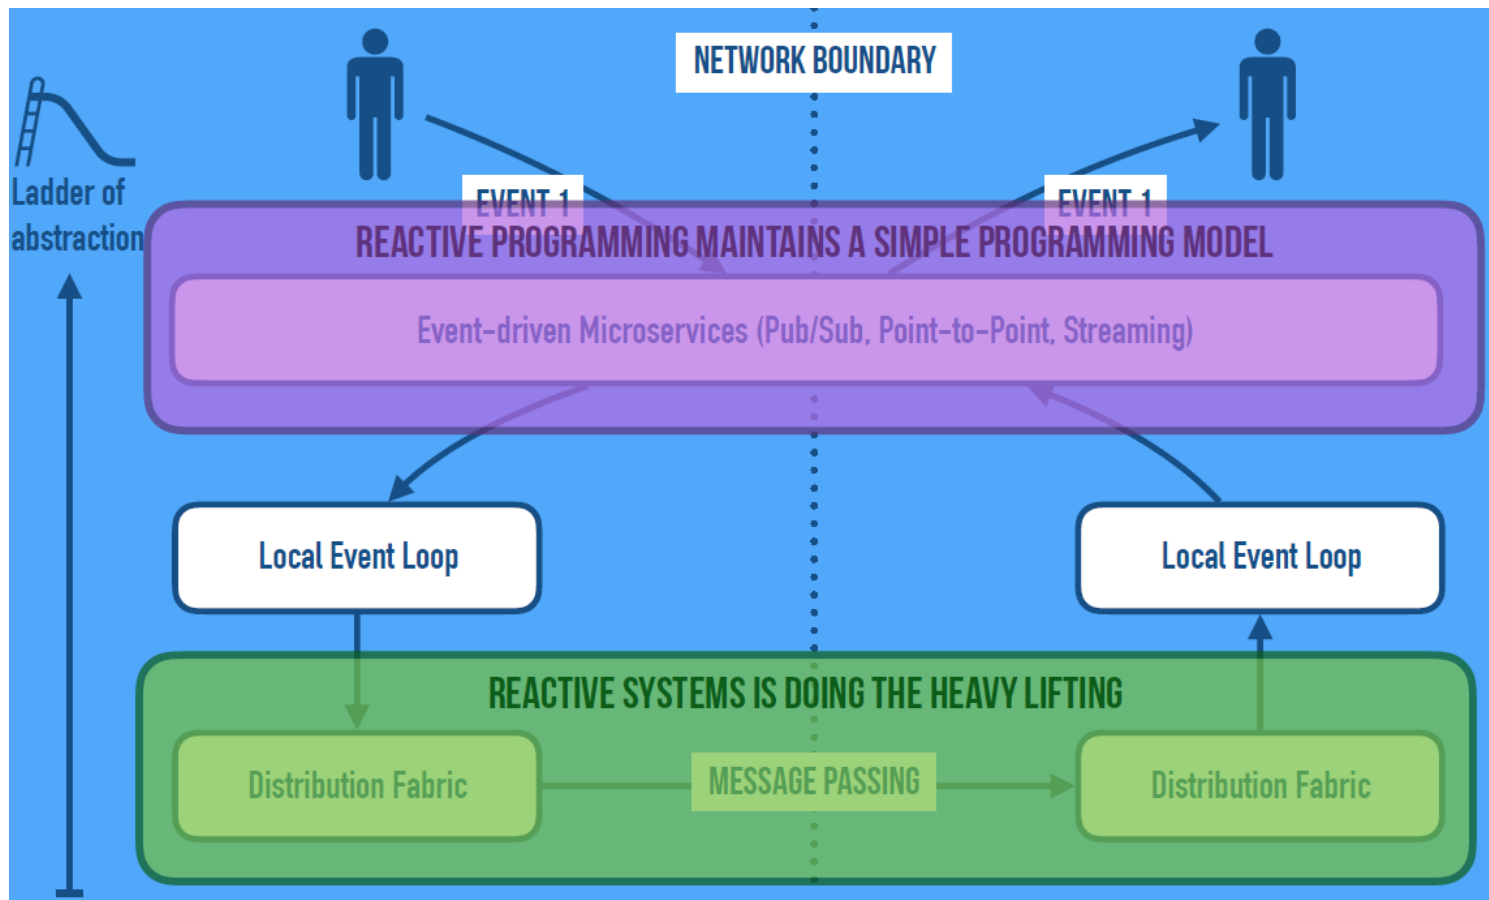


Reactive microservices and serverless computing

Internet of Things and Services

Event Driven Microservices

Event Driven Microservices Powered By Reactive Systems Fabric



Reactive microservices and serverless computing



Events First Domain Driven Design

- “When you start modeling events, it forces you to think about the behavior of the system. As opposed to thinking about the structure of the system.”
 - Greg Young, A Decade of DDD, CQRS, Event Sourcing
- Don't focus on the things (**Nouns, Domain Objects**)
- Focus on what happens (**Verbs, Events**)

Event Driven Design

* INTENTS

- ➔ Communication
- ➔ Conversations
- ➔ Expectations
- ➔ Contracts
- ➔ Control Transfer

* FACTS

- ➔ State
- ➔ History
- ➔ Causality
- ➔ Notifications
- ➔ State Transfer

COMMANDS

EVENTS

Event Driven Design

⊗ Commands

- ⊠ **Object form of** method/Action request
- ⊠ Imperative: CreateOrder, ShipProduct

⊗ Reactions

- ⊠ **Represents** side-effects

⊗ Events

- ⊠ **Represents something that** has happened
- ⊠ Past-tense: OrderCreated, ProductShipped

COMMANDS VS EVENTS

- | | |
|-----------------------------------|---|
| 1. All about intent | 1. Intentless |
| 2. Directed | 2. Anonymous |
| 3. Single addressable destination | 3. Just happens – for others (0-N) to observe |
| 4. Models personal communication | 4. Models broadcast (speakers corner) |
| 5. Distributed focus | 5. Local focus |
| 6. Command & Control | 6. Autonomy |



Event Sourcing

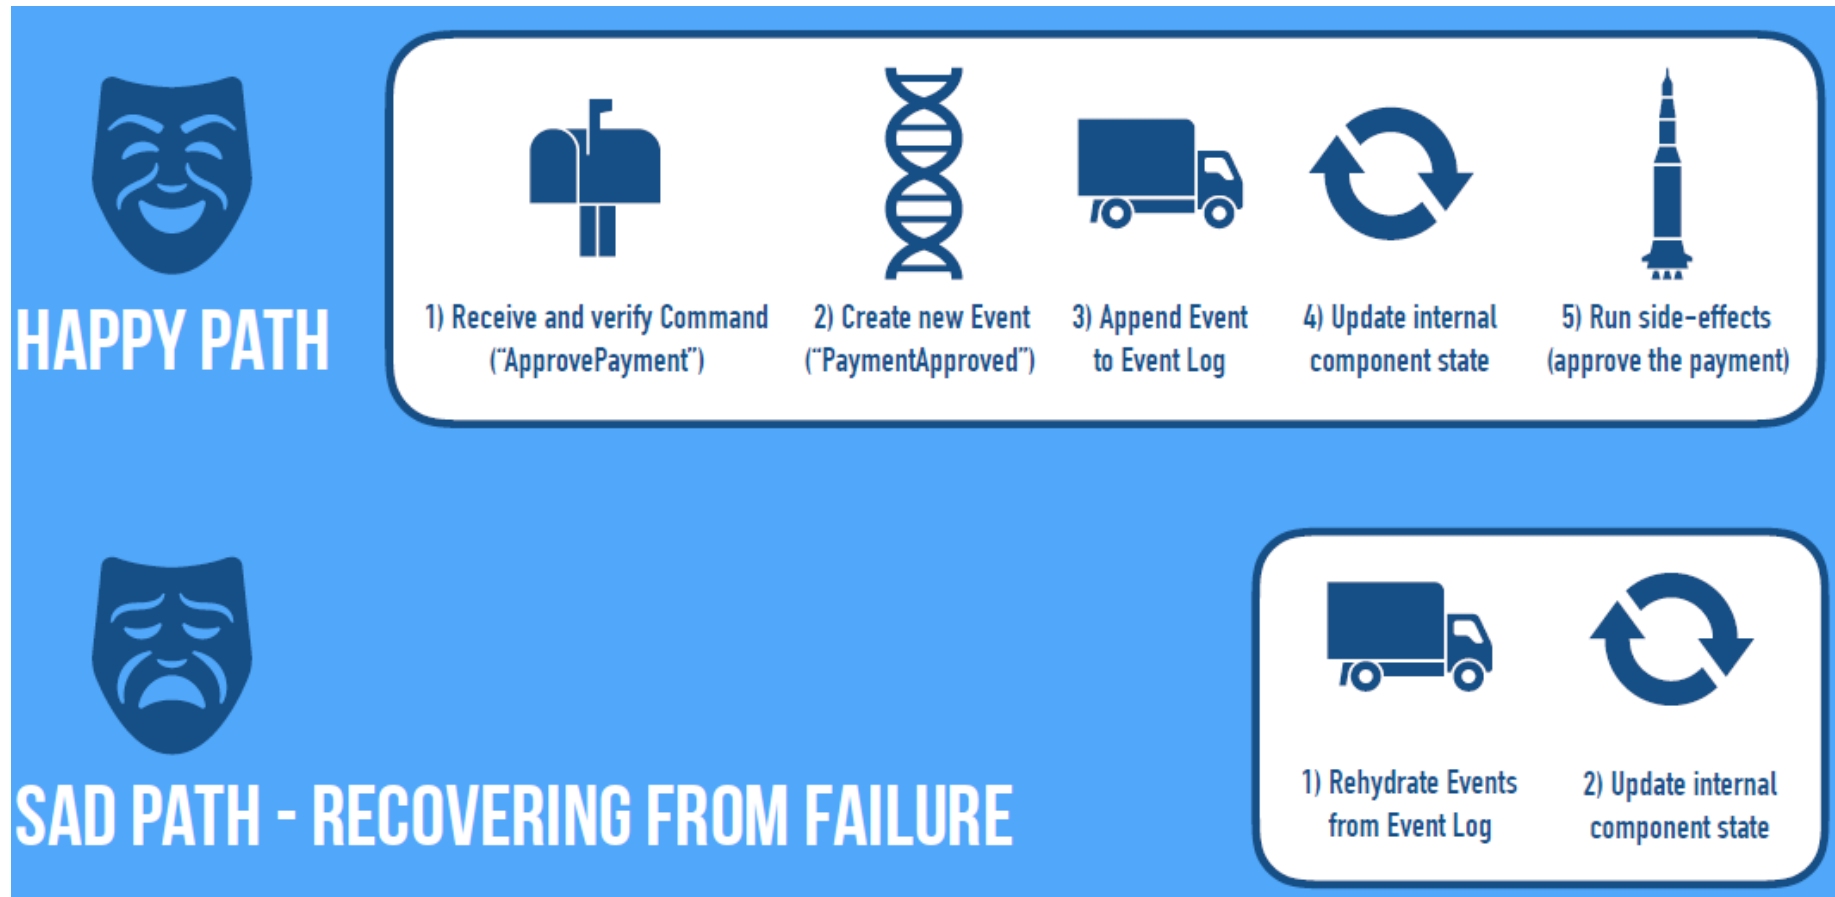
✚ Benefits

- ✚ One single **Source of Truth** with All history
- ✚ Allows for **Memory Image** - durable in-memory state
- ✚ Avoids the **Object-relational** Impedance mismatch
- ✚ Allows others to **Subscribe** to state changes
- ✚ Mechanical sympathy (**Single Writer Principle** etc.)

✚ Disadvantages

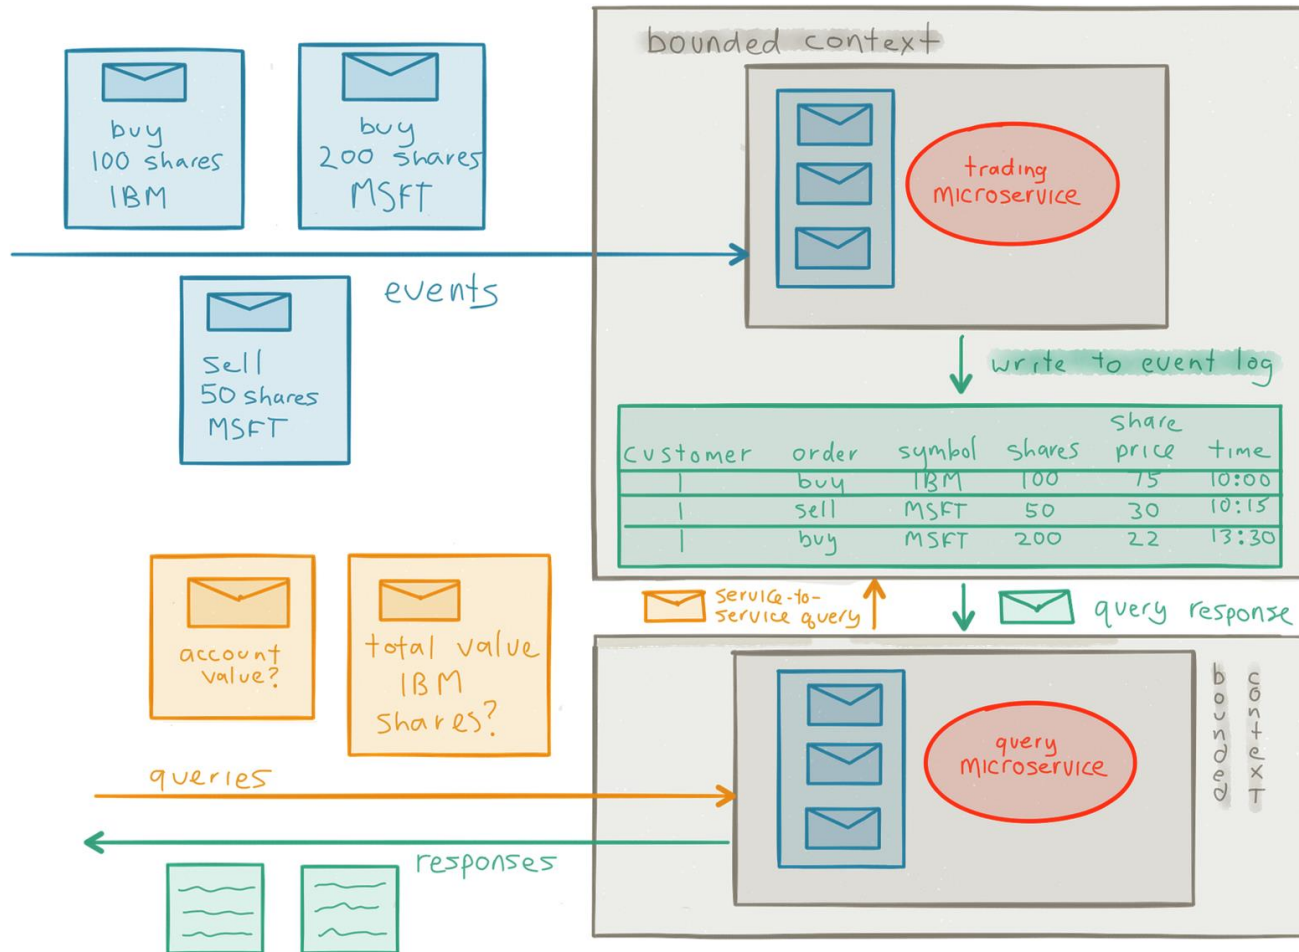
- ✚ **Unfamiliar** model
- ✚ **Versioning** of events
- ✚ **Deletion** of events (legal or moral reasons)

Event Sourced Microservices



Event Logging and CQRS

Event-based persistence through Event Logging and CQRS





Reactive Microservice Systems

✚ Key design principles

1. Don't build Microliths
2. Microservices come in (distributed) systems
3. Microservices come as (micro)systems
4. Embrace the Reactive Principles
5. Embrace Events-first Domain Driven Design
6. Embrace Event-driven Architecture
7. Embrace Event-driven Persistence



SERVERLESS COMPUTING

Serverless computing

- ❁ Cloud computing model which aims to abstract server management and low-level infrastructure decisions away from users
 - ❏ Users can develop, run and manage application code (i.e., functions), without no worry about provisioning, managing and scaling computing resources
 - ❏ Runtime environment is fully managed by Cloud provider
 - ❏ Serverless: functions still run on “servers” somewhere but we don’t care
- ❁ Function as a Service (FaaS) often as synonym

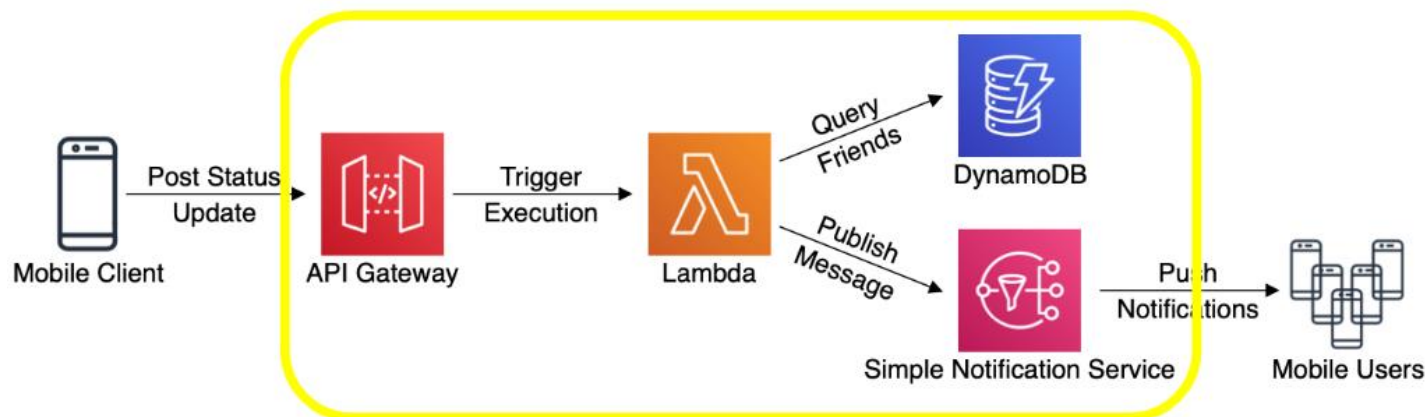
Serverless computing: characteristics

- ✿ Ephemeral compute resources
 - ✦ May only last for one function invocation
 - ✦ Automated (i.e., zero configuration) elasticity
- ✿ True pay-per-use
 - ✦ Pay only for consumed compute time, rather than on pre-purchased units of capacity
- ✿ Event-driven
 - ✦ When an event is triggered, a piece of infrastructure is allocated dynamically to execute the function code
- ✿ Can simplify the process of deploying code into production
 - ✦ Scaling, capacity planning and maintenance operations are hidden from developers or operators

Example of serverless application

Mobile backend for a social media app

1. The user composes the status update using the mobile clients of the social media platform
2. The users ends the status update using the mobile client
3. The platform orchestrates the operations needed to propagate the update inside the social media platform and to the user's friends using serverless technology
4. Each friend receives the update on their social media clients



Serverless cloud services

- ✿ Several Cloud service providers offer serverless computing on their public clouds as fully managed service
 - ✦ AWS Lambda
 - Includes lambda functions at the edge
 - ✦ Azure Functions
 - ✦ Google Cloud Functions
 - ✦ IBM Cloud Functions
- ✿ These cloud platforms also offer other supporting services (e.g., event notification, storage, database services) that are necessary for operating an overall serverless ecosystem

Example: Google Cloud Functions

❁ The “Hello World” FaaS example from Google

- ❁ HTTP request written in Node.js that displays “Hello World” or “Hello (name)” if you pass in a parameter

```
/**
 * HTTP Cloud Function.
 *
 * @param {Object} req Cloud Function request context.
 * @param {Object} res Cloud Function response context.
 */
exports.helloHttp = function helloHttp (req, res) {
  res.send(`Hello ${req.body.name || 'World'}!`);
};
```

❁ A more sophisticated example



Serverless computing: challenges and limitations

⊗ Performance

- ⊗ Startup latency and cold starts
- ⊗ “The first time you deploy a function it may take several minutes as we need to provision the underlying infrastructure to support your functions. Subsequent deployments will be much faster.” (Google Cloud Functions)

⊗ Programming language support

⊗ Resource limits

⊗ Lack of standards and vendor lock-in

Reduced flexibility



Composition of serverless functions

- ✚ Write small, simple, stateless functions
 - ✚ Complex functions are hard to understand, debug, and maintain
 - ✚ Separate code from data structures
- ✚ Then **compose** them in a **workflow**

Example: AWS Step Functions

- AWS Step Functions is a serverless orchestration service that allows the developer to coordinate multiple Lambda functions into workflows
- Example: process photo immediately after its upload in S3







Open-source FaaS platforms

- ✿ Can run on commodity hardware
- ✿ Most platforms rely on Kubernetes for orchestration and management of serverless functions
 - ✦ Configuration management of containers
 - ✦ Container scheduling and service discovery
 - ✦ Elasticity management
- ✿ Prominent platforms
 - ✦ Apache OpenWhisk - <https://openwhisk.apache.org/>
 - ✦ OpenFaaS - <https://www.openfaas.com/>
 - ✦ Nuclio - Serverless Platform for Automated Data Science - <https://nuclio.io/>
 - ✦ Fission - <https://fission.io/>
 - ✦ Knative - <https://knative.dev/>

OpenWhisk



- Open-source, distributed serverless platform that executes functions in response to events at any scale
- Based on Docker containers
- Deployed on Kubernetes, OpenShift, Mesos, Docker Compose
- Developers write functional logic (called Actions)
 - In any supported programming language
 - Can be dynamically scheduled and run in response to associated **events** (via **Triggers**) from external sources (**Feeds**) or from HTTP requests
- Functions can be combined into **compositions**




Reactive microservices and serverless computing

OpenFaaS



- Open-source FaaS framework for building functions on top of Docker and Kubernetes
- OpenFaaS stack
 - Gateway provides an external route into the functions, collects metrics and scale functions
 - Prometheus provides metrics and enables auto-scaling
 - NATS provides asynchronous execution and queuing



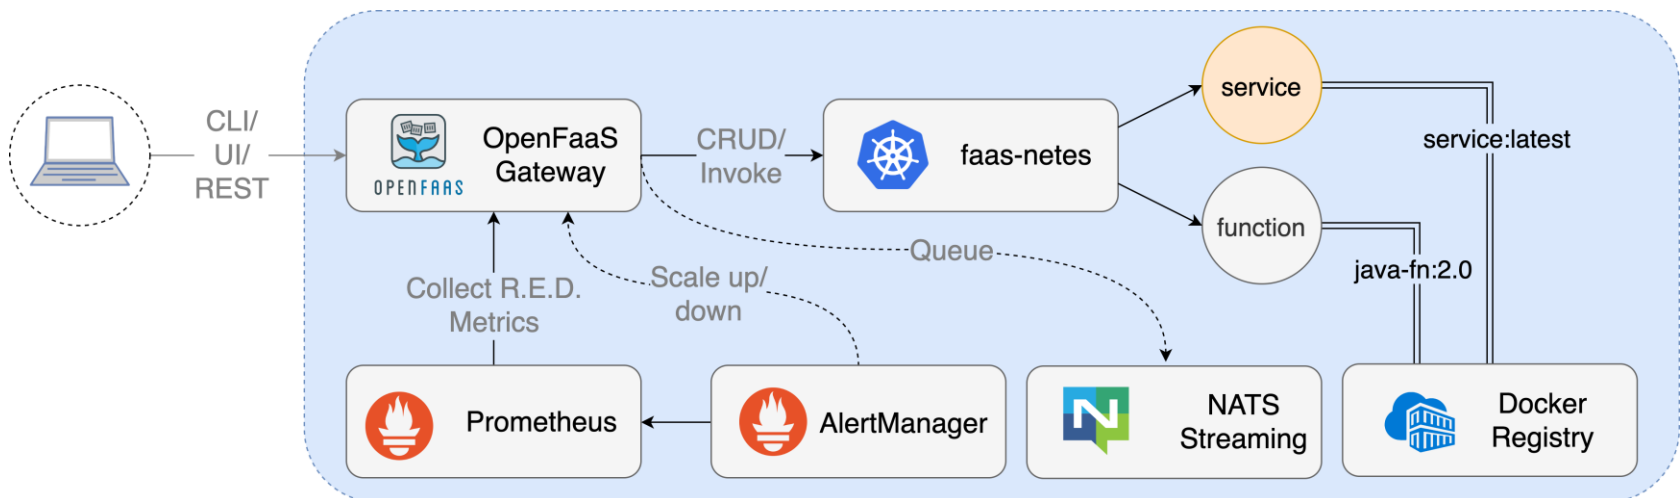
Reactive microservices and serverless computing

OpenFaaS



Conceptual workflow

- ❑ OpenFaaS Gateway can be accessed through its REST API, via CLI or through UI
- ❑ Prometheus collects metrics which are available via Gateway's API and are used for auto-scaling
- ❑ NATS Streaming enables long-running tasks or function invocations to run in the background





MICROSERVICE EXAMPLES

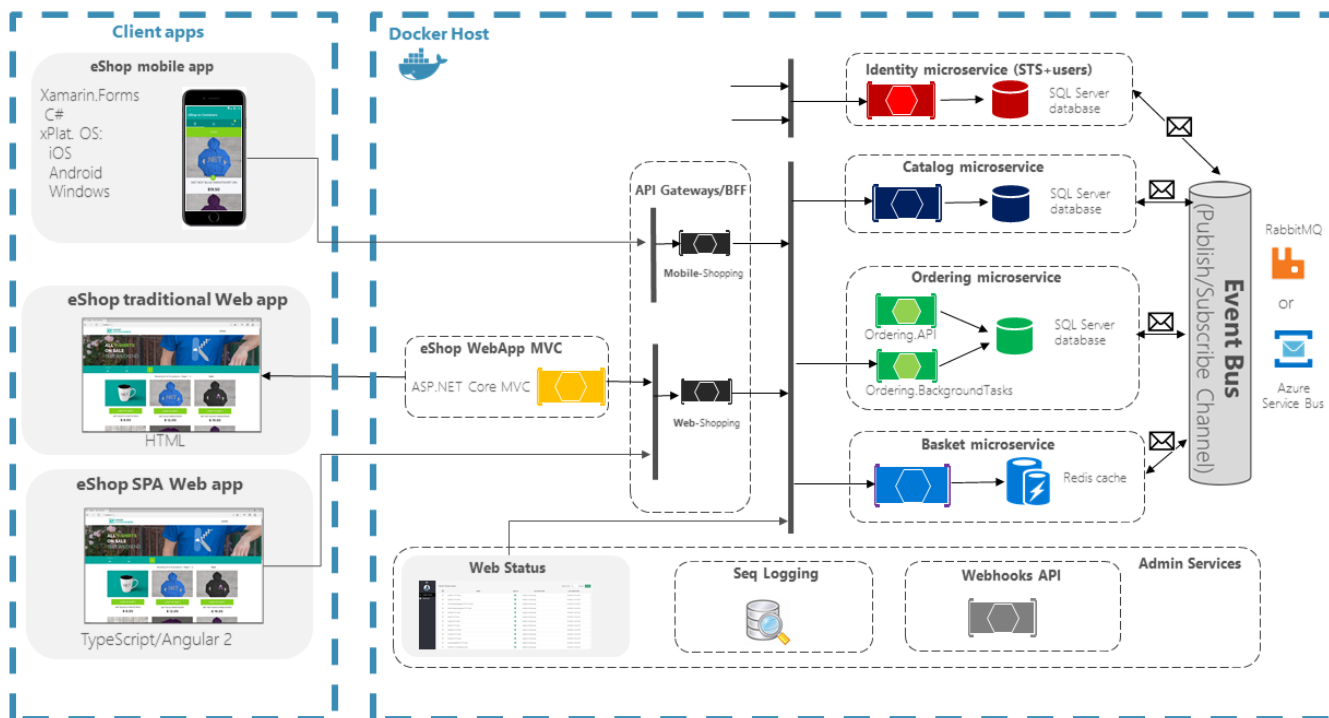
Example 1: eShopOnContainers

eShopOnContainers

- ✦ <https://github.com/dotnet-architecture/eShopOnContainers>
- ✦ <https://dotnet.microsoft.com/learn/aspnet/microservices-architecture>

eShopOnContainers reference application

(Development environment architecture)

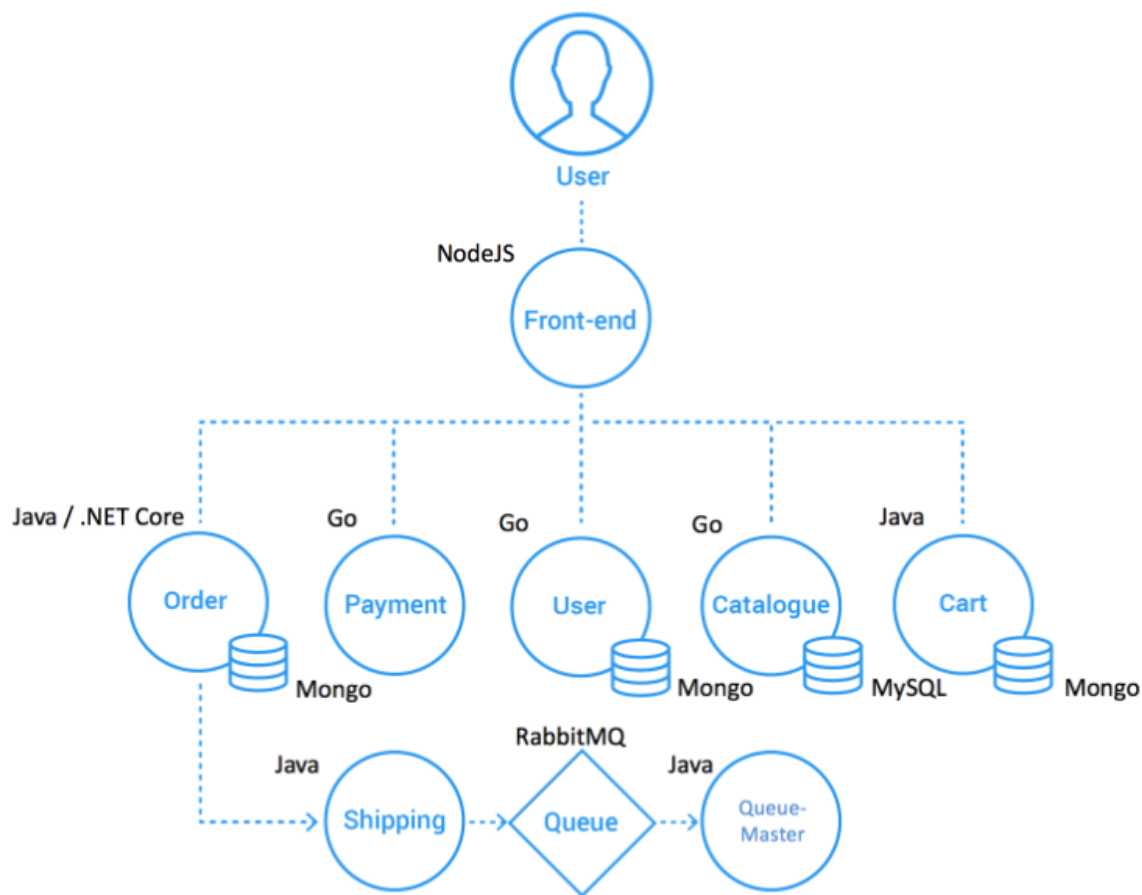


Reactive microservices and serverless computing

Internet of Things and Services

Example 2: Sock shop

- An online shop <https://microservices-demo.github.io/>

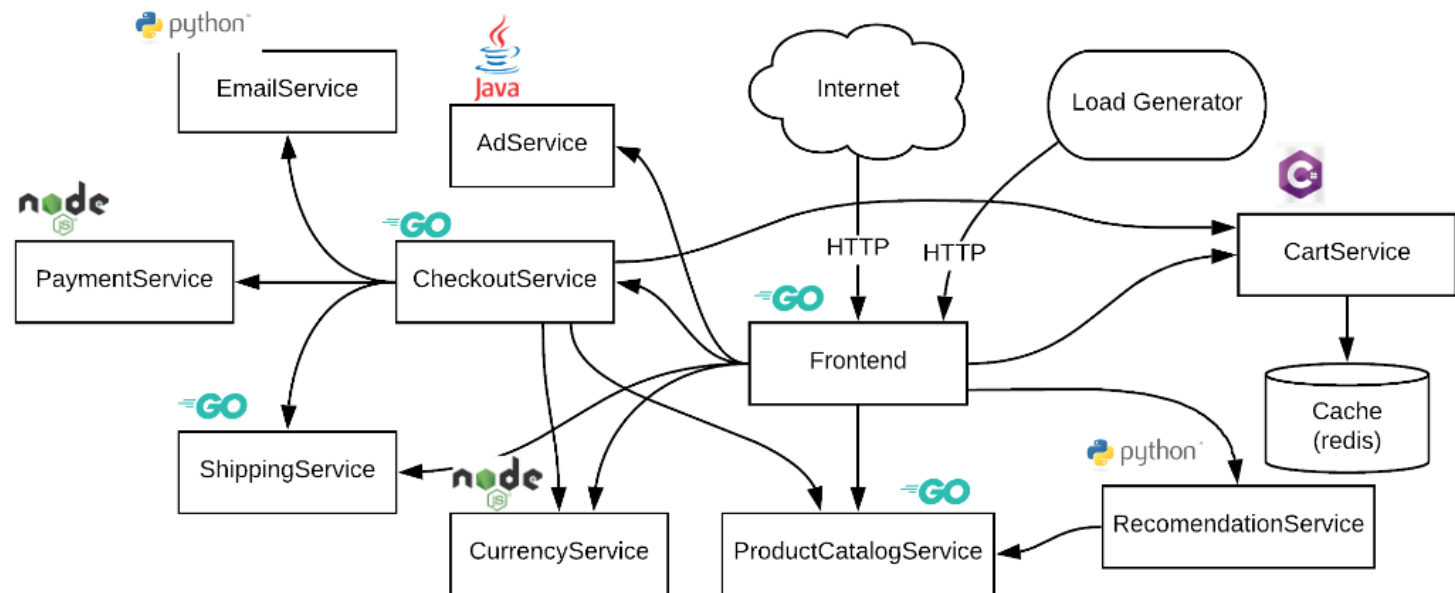


Example 3: Online boutique

Another online store

<https://github.com/GoogleCloudPlatform/microservices-demo>

- Composed of many microservices written in different languages that talk to each other over gRPC
- Used by Google to demonstrate use of Kubernetes/GKE, Istio, Stackdriver, gRPC and OpenCensus



Reactive microservices and serverless computing

Internet of Things and Services

References

- ✿ Jonas Boner, *Reactive microservice architecture: Design Principles for Distributed Systems*, O'Reilly Media Inc. 2016.
 - ✦ <https://www.lightbend.com/ebooks/reactive-microservices-architecture-design-principles-for-distributed-systems-oreilly>
- ✿ Jonas Boner, *Reactive Microsystems: The Evolution of Microservices at Scale*, O'Reilly Media Inc. 2017.
 - ✦ <https://www.lightbend.com/ebooks/reactive-microsystems-evolution-of-microservices-scalability-oreilly>
- ✿ Reactive Architecture: Foundations (CognitiveClass course)
<https://cognitiveclass.ai/learn/reactive-architecture-foundations>
 - ✦ Reactive Architecture: Introduction to Reactive Systems
 - ✦ Reactive Architecture: Domain Driven Design
 - ✦ Reactive Architecture: Reactive Microservices



References - Serverless

- ✿ M. Roberts, *Serverless Architectures*,
 - ✦ <https://martinfowler.com/articles/serverless.html>
- ✿ P. Castro et al., *The Rise of Serverless Computing*, ACM Comm., Dec. 2019.
 - ✦ <https://cacm.acm.org/magazines/2019/12/241054-the-rise-of-serverless-computing/fulltext>
- ✿ E. Jonas et al., *Cloud Programming Simplified: A Berkeley View on Serverless Computing*, Technical Report No. UCB/EECS-2019-3, 2019.
 - ✦ <https://www2.eecs.berkeley.edu/Pubs/TechRpts/2019/EECS-2019-3.pdf>